

Load Test Script

Follow Kind_Cluster Script first then do following steps :

7. Create the k6 spike-load script

Create s1_spike.js:

```
cat > s1_spike.js <<'EOF'  
import http from 'k6/http';  
import { check, sleep } from 'k6';  
import { randomSeed } from 'k6';  
randomSeed(42);  
  
export const options = {  
  summaryTrendStats: ['avg','min','med','p(90)','p(95)','max'],  
  scenarios: {  
    spike_all: {  
      executor: 'ramping-arrival-rate',  
      startRate: 10,           // warm start  
      timeUnit: '1s',  
      preAllocatedVUs: 200,  
      maxVUs: 400,  
      stages: [  
        { target: 30, duration: '60s' }, // warm up  
        { target: 400, duration: '20s' }, // sudden spike  
        { target: 400, duration: '2m' }, // hold  
        { target: 60, duration: '60s' }, // ramp down  
        { target: 0, duration: '30s' }, // drain  
      ],  
    },  
  },  
};
```

```

    },
    thresholds: {
      'http_req_failed': ['rate<0.02'], // <2% failures
      'http_req_duration{app:sock}': ['p(95)<1000'],
      'http_req_duration{app:book}': ['p(95)<1000'],
      'http_req_duration{app:tea}': ['p(95)<1500'],
    },
  };
// Default to hostPort 80 on the kind VM
const BASE = __ENV.BASE_URL || 'http://127.0.0.1';

const TARGETS = [
  { path: '/sock', tag: 'sock', weight: 0.55 },
  { path: '/book?u=normal', tag: 'book', weight: 0.20 },
  { path: '/book?u=test', tag: 'book', weight: 0.05 },
  { path: '/tea', tag: 'tea', weight: 0.20 },
];
const totalW = TARGETS.reduce((a, b) => a + b.weight, 0);

function pick() {
  let r = Math.random() * totalW;
  for (const t of TARGETS) {
    r -= t.weight;
    if (r <= 0) return t;
  }
  return TARGETS[0];
}

```

```
export default function () {
  const t = pick();
  const res = http.get(` ${BASE}${t.path}` , { tags: { app: t.tag } });
  check(res, { 'status 2xx/3xx': r => r.status >= 200 && r.status < 400 });
  sleep(Math.random() * 0.8);
}
```

EOF

Quick check:

```
ls -l s1_spike.js
head -20 s1_spike.js
```

8. Run the spike test

You can either rely on the default BASE (`http://127.0.0.1`) or set it explicitly.

```
export BASE_URL="http://localhost"
```

```
k6 run s1_spike.js
```

k6 will show:

- checks (success %)
- `http_req_duration` (avg, p95, etc.), also split by app tag
- `http_req_failed` rate

Start with lower numbers (e.g. reduce target peaks to ~150–200 req/s) if your VM is small, then increase until you find the “knee” where latency or error rates blow up.

Cluster_Load_script_updated

9. What to monitor during the spike

From the same VM:

```
kubectl top pods -A  
kubectl -n istio-system logs -f deploy/istio-ingressgateway  
kubectl -n sock-shop get pods -o wide  
kubectl -n bookinfo get pods -o wide  
kubectl -n teastore get pods -o wide
```

If you enabled the optional addons:

- **Kiali:** service graph, p95 latency, error rates.
- `kubectl -n istio-system port-forward svc/kiali 20001:20001`
- # then open `http://localhost:20001`
- **Grafana:** detailed Prometheus dashboards.
- `kubectl -n istio-system port-forward svc/grafana 3000:3000`
- # then open `http://localhost:3000`

You can also inspect traces in Jaeger (if installed) for long-tail requests during the spike.

10. Saving and comparing runs

Export results to JSON so you can compare before/after experiments (e.g. with and without HPA, different resource limits, etc.):

```
k6 run --summary-export spike_summary.json s1_spike.js
```

Repeat with different configs and compare p95 latency, error rates, and throughput over time.